

Customer & MVP Blueprint

AI-Assisted Major-Incident Pattern Detection for SMB IT Service Desks

Prepared for Azad Singh | 60 Days of Claude — Business Validation Sprint | August 17, 2026

Executive Summary

The original idea — a generic AI triage layer for SMB IT service desks — is real pain but a populated market: eesel AI already sells an almost identical wedge, and Freshservice, Jira Service Management and Zendesk are closing the gap natively. This blueprint carries forward the validated PIVOT: narrow the product to Major-Incident Pattern Detection layered on top of standard triage — using Azad's 12+ years of major incident management (MIM) experience to spot ticket clusters that signal a brewing major incident before a human notices. This is a genuinely underserved angle none of the researched competitors emphasize, and it converts Azad's strongest personal credibility into the product's core defensibility.

Recommended path: do not build yet. Run the 30-day discovery-and-fake-door plan already defined in the validation report, using this blueprint's ICP, personas, objections and MVP scope as the script for those conversations.

Ideal Customer Profile

Attribute	Definition
Company size	50–500 employees (sweet spot 100–300)
IT team size	2–8 person IT/service desk team, no dedicated automation or AIOps role
Current tooling	Freshservice (Growth/Pro tier) or Jira Service Management — actively on a modern cloud ITSM
Ticket volume	300–3,000 tickets/month
Industry fit	Vertical-agnostic start; strongest early fits: professional services, healthcare admin, education, logistics
Budget	\$100–500/month realistic bolt-on spend; below ~\$6K/year usually skips formal procurement
Buying trigger	Ticket volume outpacing headcount, a bad SLA-miss quarter, or a new IT Manager's first-90-days modernization push

Buyer Persona

Field	Detail
Title	IT Manager / Service Desk Lead / IT Director wearing an ops hat
Owns	The SLA dashboard; personally fields “why is this ticket still unassigned” escalations
Vocabulary	Priority matrices, category taxonomies, MTTR, SLA tiers, major incident process
Fear	Missing the early signal of a brewing major incident until it's already a P1 with executives asking questions
Evaluation habit	Has likely already looked at Freddy AI, eesel AI, or Rovo and either rejected the price tier or found it “just categorization, not real intelligence”
Trust driver	Wants auditable, explainable suggestions — not a black-box “agent” taking autonomous action

Top 10 Customer Pain Points

- Manually re-reading ticket streams to spot patterns before they escalate into a major incident
- AI triage capability exists but is gated behind pricing tiers they haven't bought
- No dedicated automation/AIOps headcount to build detection logic in-house
- SLA misses discovered only after several related tickets have already piled up
- Fatigue evaluating “agentwashed” AI tools that overpromise and underdeliver
- Fear of black-box AI making triage or escalation calls without a clear rationale
- Cross-referencing tickets against the knowledge base is slow and inconsistent across agents
- Junior agents miscategorize or under-prioritize tickets that are early major-incident signals
- Leadership only hears about a brewing incident once it is already visible and painful
- Switching ITSM platforms to get better AI is too disruptive to justify

Customer Journey

Awareness → Consideration → Purchase → Retention

Stage	What happens
Awareness	Hears of the tool via practitioner-led content (LinkedIn, ITSM/sysadmin communities), a referral from an ex-colleague, or after a bad SLA-miss postmortem prompts a search for “major incident detection”
Consideration	Compares against eesel AI / native Freddy-Rovo triage; evaluates whether pattern-detection is meaningfully different from “just categorization”; asks for a concierge or pilot trial on real tickets
Purchase	Low-friction, no-RFP purchase under ~\$6K/year; decision driven by the IT Manager alone after seeing real output quality on their own ticket history
Retention	Stays because detected clusters actually pre-empted a major incident at least once; churns if the tool is noisy (false positives) or indistinguishable from native triage

Key Customer Objections

- “We already have Freddy AI / Rovo triage — how is this different?”
- “We can't afford another tool on top of what we already pay for.”
- “Who built this, and why should we trust a solo/early-stage vendor with our ticket data?”
- “Will this create alert fatigue with false-positive major-incident flags?”
- “Can it plug into our exact ITSM tool without migration or heavy setup?”

Key Buying Triggers

- A recent major incident that, in hindsight, showed early warning signs across multiple tickets
- A new IT Manager's first-90-days initiative to modernize the desk
- Ticket volume growth outpacing headcount
- Rejection of a native AI add-on due to pricing-tier gatekeeping

MVP Recommendation

What should be built first

- A cluster-detection layer that reads new and recent tickets from one ITSM tool (Freshservice or Jira Service Management) and flags patterns resembling the early stage of a major incident

- A simple, auditable explanation attached to every flag — which tickets, what pattern, why it was raised — no black-box scoring
- A concierge-mode delivery for the first 3–5 design partners: Azad manually runs the detection using an LLM plus the customer's own ticket history and knowledge base, before any product exists

What should NOT be built

- Generic per-ticket triage (category/priority/first-fix suggestion) — that wedge is already owned by eesel AI and native vendors
- Autonomous resolution or auto-actioning tickets — trust in this ICP is won by transparency, not automation
- Multi-platform support (Zendesk, ServiceNow, SysAid) before the Freshservice/JSM wedge is proven

Success metrics

- At least 3 of 5 concierge customers report the detected pattern would have changed how early they escalated
- Landing page visit-to-signup conversion from the fake-door test clears a meaningful double-digit rate among warm, pain-qualified traffic
- 3–5 design partners convert to paid pilots after the concierge phase

MoSCoW Prioritization

Priority	Items
Must Have	Ticket ingestion from one ITSM platform; cluster/pattern detection logic; auditable explanation per flag; concierge delivery workflow
Should Have	Simple dashboard or Slack/Teams alert for flagged clusters; feedback loop (was this flag useful?) to tune detection
Could Have	Second ITSM platform integration; KB-grounded suggested first response per flagged cluster
Won't Have (v1)	Autonomous ticket resolution; multi-platform day-one support; enterprise/ServiceNow tier

Pricing Hypothesis

Anchor near the validated \$100–500/month bolt-on budget band, priced per organization rather than per agent to stay simple for a 2–8 person desk. A usage-based option (per-ticket-analyzed, similar to eesel's per-ticket model) may reduce purchase friction for skeptical first-time buyers. Concierge/pilot customers should be priced low or free in exchange for detailed feedback and a case-study reference.

Top 5 Risks

- Native vendors (Freshservice, Atlassian) extend triage into pattern/cluster detection, closing the wedge from the inside
- Detection logic produces false positives, causing alert fatigue and churn
- Founder has no AI/ML engineering background — building a trustworthy detection pipeline may require a technical co-founder
- Divided founder attention across a full-time role, two other career paths, and a job search limits sustained execution
- Discovery calls reveal the pivot angle isn't meaningfully differentiated from generic triage in customers' minds

30-Day MVP Plan

Week	Focus
Week 1	Write a 10-question discovery script; build a 60–80 person target list of IT Managers on Freshservice/JSM; draft practitioner-led outreach

Week	Focus
Week 2	Run 10–15 discovery calls testing pain around missed early major-incident signals and reactions to the pattern-detection pivot
Week 3	Finish remaining calls; synthesize whether the pivot gets a materially stronger reaction than generic triage; identify 3–5 unprompted “I’d pay for this” leads
Week 4	Publish a narrowed fake-door landing page; track conversion; run a free concierge MVP for top leads using Claude plus their own KB; decide Go / No-Go / Pivot-again

Founder Action Sheet — Top 10 Next Actions

- Draft the 10-question discovery script this week
- Build the 60–80 person IT Manager target list from personal network and ITSM communities
- Write practitioner-led outreach copy that leads with MIM credibility, not a pitch
- Book 10–15 discovery calls before writing any code
- Explicitly ask each call whether they've evaluated eesel AI, Freddy AI, or Rovo and why they did or didn't adopt
- Log every call in a pain-score spreadsheet (1–5, current tool, AI-add-on awareness, willingness to pay)
- Identify and shortlist 3–5 unprompted “I’d pay for this” design-partner candidates
- Build a single fake-door landing page for the narrowed pivot pitch
- Run a manual concierge MVP for top leads using Claude/an LLM plus their ticket history and KB
- At Day 30, decide Go / No-Go / Pivot-again strictly from call signal, landing-page conversion, and concierge feedback

Scores (0–100)

Dimension	Score	Basis
Customer Clarity	80	ICP and persona are well-defined and corrected against real tooling patterns
Problem Severity	60	Pain is real but partially already addressed by existing native and third-party tools
PMF Potential	45	Depends entirely on whether the pattern-detection wedge tests as meaningfully different in real conversations
MVP Readiness	35	No technical co-founder or ML engineering capability confirmed yet; concierge-first approach mitigates this

Final Verdict

Promising but Unvalidated

The core pain (missed early major-incident signals) is credible and tied to genuine founder expertise, but it has not been tested against real IT Managers, and the broader “generic AI triage” framing is already a populated, competitive market. Proceed to the 30-day discovery plan before any build decision.